

Reimagine CRM around the buyer’s decision

DecisionOS

An AI-native decision layer for complex enterprise sales.

INITIAL WEDGE · COMPLEX ENTERPRISE B2B NEW-LOGO SALES

Qualified opportunity → commercial decision

CRM TODAY

Seller process

Opportunity

Stage

Activities

Forecast

DECISIONOS

Buyer decision

Stakeholders

Criteria

Commitments

Decision gaps

“What did we do?”

“What must become true?”

Traditional CRM models the seller’s process. DecisionOS models the buyer’s decision.

MARKET SHIFT

The buyer is more autonomous — and the decision more complex

That makes human seller attention less frequent, but more consequential.

67%

prefer a rep-free
buying experience

69%

prefer validating AI-generated
insight with a salesperson

13 + 9

typical buying decision: 13 internal
stakeholders + 9 external influencers

MORE SELF-SERVICE**Buyer research shifts to digital + AI**

Sellers are no longer the primary information source at every step.

MORE COMPLEXITY**Finance · Security · Procurement · IT**


Different stakeholders carry different risks, incentives and veto power.

AI makes sales execution cheaper. Human attention becomes more valuable.

COMPETITIVE REALITY

Decision intelligence is emerging. What becomes the primary object?

Buyer-side signals, stakeholder gaps and next actions already exist. Our bet is to make decision state the primary object — not another signal feeding an opportunity.



We are not claiming buyer-decision thinking is new. The differentiation is a temporal, evidence-backed decision state as the operating object.

THE BROKEN ABSTRACTION

CRM stage ≠ buyer progress

A seller can move the opportunity forward while the buying group remains unresolved.

WHAT CRM SEES

Acme Corp

\$240K ARR · Proposal · Nov 20

- ✓ Champion engaged
- ✓ Technical demo complete
- ✓ Security started
- ✓ Proposal sent

Looks healthy

WHAT ACTUALLY CHANGED

“Finance wants to understand whether the implementation effort is justified. They will probably want payback inside 12 months.”

Sarah · Champion · Oct 2

Problem	VALIDATED
Technical	VALIDATED
Security	IN PROGRESS
Economics	UNRESOLVED
Procurement	UNKNOWN
Consensus	INCOMPLETE

The opportunity progressed to proposal. The buyer’s economic decision did not.

SCOPE ♦ ICP

Scope: complex enterprise B2B new-logo sales

Qualified opportunity → commercial decision. Narrow wedge first; CRM replacement is explicitly out of scope.

PRIMARY USER

Enterprise Account Executive

\$50K♦

ACV

60♦ days

sales cycle

5♦

stakeholders

8–25

active deals

JOB TO BE DONE

“When I manage multiple complex opportunities, tell me what is actually preventing each buying group from reaching a decision — and where my involvement can materially change the outcome.”

SECONDARY USER

Front-line Sales Manager

Needs intervention visibility, not another dashboard.

PRODUCT BUYER

CRO / VP Sales / RevOps

Buys better revenue execution and seller capacity.

IN SCOPE

New-logo opportunity · 5+ stakeholders · Finance / Security / IT / Procurement


OUT OF SCOPE

B2C · transactional SMB · AI SDR / lead gen · renewals / support · CRM replacement

NEW PRODUCT PRIMITIVE

The Decision Twin: a live model of what the buyer needs to decide

Persistent structured state — not an ephemeral ITM summary.




Current state → Decision gap → Intervention → Human / AI action → Buyer response → Updated state

PROTOTYPE

Don't give sellers more tasks. Give them fewer, higher-leverage decisions

Homepage = exception queue for moments where human judgment matters.



WHAT THE UX PROVES

18 deals → 4 human moments

14 need no human action now; AI monitors and prepares instead of creating more alerts.

Two decision-gap types

Acme = Finance validation. Northstar = stalled Security approval.

AI works before asking

Brief, ROI worksheet and intro draft are already prepared.

Chat is secondary


The seller should not need to know what question to ask.

CORE WORKFLOW

Evidence → intervention → buyer state change


The product tracks whether the decision moved — not whether the seller completed a task.

1 · UNDERSTAND




Decision Twin identifies Economic validation gap

2 · VERIFY



Facts, reports, inferences and unknowns stay distinct

3 · ACT



AI prepares; seller owns relationship + judgment

Second full prototype path: Northstar Bank · missed Security commitment → blocker discovery → explicit approval criterion.

COLD-START MVP

Valuable before it learns

No proprietary historical outcome data is required for version one.

1

EVIDENCE

CRM · email · calendar · meeting transcripts

2

DECISION ONTOLOGY

Stakeholders · criteria · approvals · commitments

3

EXPLICIT RULES

Explainable gap logic for common enterprise decisions

4

LLM REASONING

Interpret ambiguity · propose interventions · abstain when unknown

BUILD NOW

- ✓ CRM overlay
- ✓ Decision Twin
- ✓ Gap detection
- ✓ Evidence + confidence
- ✓ Intervention recs
- ✓ AI preparation
- ✓ Human approval

NOT YET


- Predictive win score
- Autonomous negotiation
- Full CRM replacement
- AI SDR
- Self-learning optimization
- Causal impact claims

First reconstruct reality. Then learn what changes outcomes.

ENGINEERING

Probabilistic intelligence, deterministic control

Use AI where interpretation adds value; use software where correctness matters.



GUARDRAILS

Provenance

Every factual claim points to source evidence.

Abstention

Unknown beats a plausible hallucination.

Risk-tiered autonomy

Internal/reversible work can automate; external action requires approval initially.

Untrusted content

Emails and transcripts are data — never instructions.

Auditability

Observe → infer → recommend → act → outcome.

Engineering principle **Deterministic where correctness matters · Probabilistic where interpretation helps · Human-controlled where consequences are high**

DELIVERY PLAN

Validate trust before autonomy — and behavior before revenue claims

Start in shadow mode; earn the right to automate.



Synthetic scenarios test system behavior. Real design-partner pilots test business value.

BUSINESS + LONG-TERM ADVANTAGE

Start as an intelligence layer. Become the system that learns how customers decide.

Sell measurable revenue execution — not model tokens.

MVP WEDGE

Decision intelligence over existing CRM

Product buyer: CRO / VP Sales / RevOps

Model: enterprise platform fee + seller
seats

No forced Salesforce migration.

LEARNING LOOP



DEFENSIBILITY

- ✓ **Decision ontology**
- ✓ **Decision histories**
- ✓ **Human corrections**
- ✓ **Intervention → outcome data**

The moat is not the LLM.

CRM learned to record what sellers did. AI CRM learned to help sellers act. The next CRM should understand how customers decide.

DecisionOS · Automate work. Escalate judgment.

APPENDIX A

Competitive pressure test: pieces already exist

The claim is not “nobody tracks buyers.” The bet is how the product is organized, governed and learns over time.

PRODUCT	WHAT EXISTS	PRIMARY OPERATING MODEL
Salesforce	Prospecting, engagement, meeting prep, pipeline updates, quoting	Opportunity + CRM context
Dynamics 365	System of action, autonomous agents, Next Best Action	Cross-agent action prioritization
HubSpot	Self-updating CRM, Growth Context, agents	Unified context + outcomes
Gong	Revenue graph, custom agents, governance, conversation intelligence	Revenue context + execution
CopilotGTM	Buyer progress, stakeholder gaps, qualification gaps, risks, next actions	Deal execution / buyer progress
MEDDPICC	Economic buyer, decision criteria/process, champion, paper process	Qualification methodology
DecisionOS	Temporal Decision Twin → evidence → gap → intervention → human/AI action → outcome	Buyer decision state + learning loop

Sources: Salesforce, Microsoft, HubSpot, Gong, CopilotGTM and MEDDPICC product/methodology pages reviewed Oct 2026.

APPENDIX B

Decision Twin schema

Structured, temporal and provenance-aware.

Opportunity	CRM metadata; seller-owned record
Stakeholder	Person, role, influence, relationship
Expressed stance / concern	Source-linked support, concern or uncertainty
Decision criterion	Condition the buyer needs satisfied
Approval requirement	Security, Finance, Legal, Procurement...
Commitment	Who promised what, by when

Evidence	Canonical source pointer + excerpt
Inference	Model interpretation + confidence
Decision gap	Current state vs required state
Intervention	Proposed change to resolve a gap
Outcome	Buyer response / state transition

Temporal rule: Never overwrite the old decision state. Store the transition so “what changed?” remains answerable.

APPENDIX C

Risk-tiered autonomy

Autonomy is a policy decision, not a model personality.

0

OBSERVE

Read, extract, summarize

Automatic

1

INTERNAL + REVERSIBLE

Refresh Decision Twin, prepare briefs

Automatic

2

ENTERPRISE WRITES

CRM notes, tasks, reminders

Permissioned

3

CUSTOMER-FACING

Email, meeting scheduling

Human approval in MVP

4

HIGH CONSEQUENCE

Pricing, contract, negotiation, legal

Human-owned

Example: AI can draft the Finance introduction. Maya must approve the customer-facing send. Pricing commitments remain human-owned.

APPENDIX D

Evaluation framework

“LLM accuracy” is not one metric. Test every layer and the full workflow.

- 1

Extraction

Stakeholders, criteria, commitments
- 2

Grounding

Claim ↔ supporting evidence
- 3

Decision state

Known / unknown / conflict / resolved
- 4

Gap detection

Important missing condition
- 5

Recommendation

Logical, feasible, safe
- 6

Action

Correct tool + params + permissions
- 7

End to end

Signal → state → gap → action → outcome

- Precision / recall
- Unsupported claim rate
- Human-labeled correctness
- False positive / false negative
- Expert rubric + accept/modify/dismiss
- Tool/action compliance
- Golden scenario pass rate

No real dataset yet?

Create a golden synthetic scenario set to test behavior — then move to shadow-mode evaluation on permissioned real deals.

APPENDIX E

Illustrative design-partner pilot

A ~12-week pilot can test the core product risk if integration scope stays deliberately narrow.

W1-2	Discovery + ontology	Design partner interviews; golden eval set; choose one CRM + one mail provider
W3-4	Evidence layer	Auth, sync, event model, source provenance, permissions
W5-6	Decision Twin	Extraction, entity resolution, temporal state, correction UX
W7-8	Gap engine	Explicit rules; conflict handling; confidence + unknown states
W9-10	Intervention	Attention Queue; planner; brief/draft generation; approval policy
W11-12	Pilot hardening	Offline evals; prompt-injection tests; shadow pilot; instrumentation

Pilot scope assumption: One CRM + one email provider + existing meeting transcripts. Production hardening and broader integrations follow only if the core hypotheses pass.

APPENDIX F

Sources, assumptions & integrity notes

Current-market research supports the problem framing; prototype data and product claims are explicitly synthetic.

Gartner · May 20, 2026	69% validate AI insights with reps; 67% rep-free; 70% digital self-service.	gartner.com/en/newsroom/press-releases/2026-05-20-...
Forrester · Feb 23, 2026	73% of purchases span 3+ departments; avg. 13 internal + 9 external influencers.	forrester.com/blogs/three-realities-about-b2b-buying-networks
Microsoft Dynamics 365 · 2026 Wave 1	“CRM from a system of record into a system of action”; agents + prioritized actions.	learn.microsoft.com/en-us/dynamics365/release-plan/2026wave1/sales/...
Salesforce · Feb/Mar 2026	87% of sales orgs use AI; Agentforce Sales covers prospecting → quoting.	salesforce.com/news/stories/state-of-sales-report-announcement-2026
HubSpot · Sep 16, 2026	Self-updating Smart CRM + Growth Context + Breeze Assistant.	hubspot.com/company-news/fall-26-spotlight
Gong · Jun 24, 2026	Revenue Harness: governed custom agents with context, permissions and oversight.	gong.io/press/gong-launches-mission-big-dipper-revenue-harness
CopilotGTM + MEDDPICC · current	CopilotGTM tracks buyer progress/stakeholder gaps; MEDDPICC formalizes economic buyer, decision criteria/process and champion.	copilotgtm.com · meddicc.com

Integrity note: Acme Corp, Maya, all deal values, prototype interactions and future product performance are fictional. No win-rate lift, prediction accuracy or proprietary data claim is made.